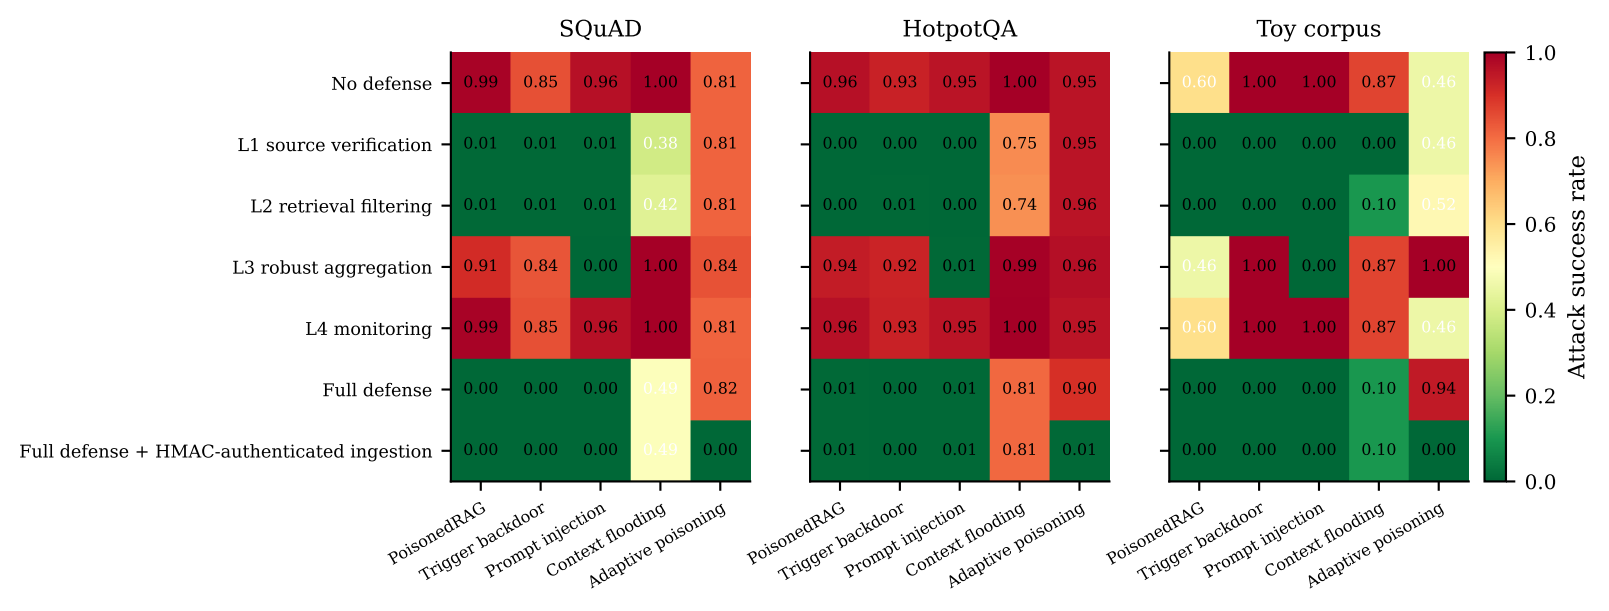
SQuAD
HotpotQA
Toy corpus
No defense
L1 source verification
L2 retrieval filtering
L3 robust aggregation
L4 monitoring
Full defense
Full defense + HMAC-authenticated ingestion
0.99 0.85 0.96 1.00 0.81
0.01 0.01 0.01 0.38 0.81
0.01 0.01 0.01 0.42 0.81
0.91 0.84 0.00 1.00 0.84
0.99 0.85 0.96 1.00 0.81
0.00 0.00 0.00 0.49 0.82
0.00 0.00 0.00 0.49 0.00
0.96 0.93 0.95 1.00 0.95
0.00 0.00 0.00 0.75 0.95
0.00 0.01 0.00 0.74 0.96
0.94 0.92 0.01 0.99 0.96
0.96 0.93 0.95 1.00 0.95
0.01 0.00 0.01 0.81 0.90
0.01 0.00 0.01 0.81 0.01
0.60 1.00 1.00 0.87 0.46
0.00 0.00 0.00 0.00 0.46
0.00 0.00 0.00 0.10 0.52
0.46 1.00 0.00 0.87 1.00
0.60 1.00 1.00 0.87 0.46
0.00 0.00 0.00 0.10 0.94
0.00 0.00 0.00 0.10 0.00
PoisonedRAG
Trigger backdoor
Prompt injection
Context flooding
Adaptive poisoning
Attack success rate
1.0
0.8
0.6
0.4
0.2
0.0